

Question 1

Part1:

For part one I tried different LR's of 0.1, 0.01, 0.001. 0.1 and 0.01, did not perform well at all with any of the three optimizers, Adam, SGD, and RMSProp. However I did manage to get the model to train at 0.001, with batch size at 128 and 25 epoches.

Adam

accuracy: 0.9965

Test loss: 0.057167619466781616

Test accuracy: 0.9882000088691711

SGD:

accuracy: 0.9930

Test loss: 0.048220593482255936

Test accuracy: 0.9846000075340271

RMSProp:

accuracy: 0.9943

Test loss: 0.08929948508739471

Test accuracy: 0.9871000051498413

Since the accuracy was at 99% I did not explore different batch sizes, also since the model was not increasing in accuracy from the 20th epoch to the 25th epoch, I didn't change the epoches to more epoches.

Part 1.2

Part 2 had very similar results, and again it seems like the batch size and amount of epoches I have picked worked out just fine.

Adam

accuracy: 0.9949

Test loss: 0.038366544991731644

Test accuracy: 0.9883000254631042

SGD: 0.9856

Test loss: 0.04696971923112869

Test accuracy: 0.984499990940094

RMSProp:

accuracy: 0.9918

Test loss: 0.06024818867444992

Test accuracy: 0.9876999855041504

Part 1.3

Part 3 has the same conclusion as part 1 and part 2.

Adam:

accuracy: 0.9952

Test loss: 0.04177796095609665

Test accuracy: 0.9890000224113464

SGD:

accuracy: 0.9844

Test loss: 0.08808410912752151

Test accuracy: 0.9700000286102295

RMSProp:

accuracy: 0.9946

Test loss: 0.06202826648950577

Test accuracy: 0.9878000020980835

Question 2

Part2.1: What is the effect of learning rate on the training process? Which performed best?

A LR of 0.01 was too much of a learning rate for the model to train with results of:

accuracy : 0.2486

Test loss: 129.45310974121094

Test accuracy: 0.2443999946117401

(note: The parameters for this is, batch size of 128 and 100 epoches, also I only used adam for this part since it seemed to perform better when testing at lr of 0.001.)

A lr of 0.001 seemed to do well, at 25 epoches and batch size of 128:

Adam: 0.5887

Test loss: 1.2418910264968872

Test accuracy: 0.5742999911308289

SGD: 0.1000

Test loss: nan

Test accuracy: 0.10000000149011612

RMSProp:0.5705

Test loss: 1.2877159118652344

Test accuracy: 0.5509999990463257

100 epoches, batch size of 128, lr 0.001, using adam:

accuracy: 0.6459

Test loss: 1.1973625421524048

Test accuracy: 0.598800003528595

1000 epoches, batch size of 128, lr 0.001, using adam:

accuracy: 0.6702

Test loss: 1.029354245464565

Test accuracy: 0.6225636562454

Any LR lower than 0.001 seemed to take too long to train, as it was not learning fast enough.

**Part2.2: What is the effect of batch size on the training process? Which performed best?
A smaller batch size takes longer to train. The bigger the batch size here the better the accuracy.**

Lr 0.001, adam, 100 epoches:

Batch size 2:

accuracy: 0.3634

Test loss: 1.9981257915496826

Test accuracy: 0.3172999918460846

Batch size 4:

accuracy: 0.4553

Test loss: 1.6594843864440918

Test accuracy: 0.433499991893768

Batch size 8:

accuracy: 0.5266

Test loss: 1.51930570602417

Test accuracy: 0.4875999987125397

Batch size 16:

accuracy: 0.5771

Test loss: 1.2947492599487305

Test accuracy: 0.5590999722480774

Batch size 32:

accuracy: 0.6125

Test loss: 1.2311395406723022

Test accuracy: 0.5827999711036682

Batch size 64:
accuracy: 0.6345
Test loss: 1.2209553718566895
Test accuracy: 0.5906000137329102

Batch size 128:
accuracy: 0.6445
Test loss: 1.1978529691696167
Test accuracy: 0.5909000039100647

Part2.3

accuracy: 0.92%
train loss: 0.038
Test accuracy: 70%
Parameters are in note book

Part2.4

accuracy: 0.7031
Test loss: 2.2949001789093018
Test accuracy: 0.4402000010013580

LeNet:
(cov =)
Conv1: 1520
Conv2: 20040
Conv3: 60060
FC1: $(300*84)+1*120 = 5124$
FC2: $(84*10)+1*10 = 850$
Total params: 87,594

Fully connected:
Fc1: 24
Fc2: 112
Fc3: 2040
Fc4: 645204
Fc5: 850
Total params: 648,230
(got with `print(model.summary())`)

No, the model started to over fit and the test accuracy is worse.

Question3:

Part3.1: What are the dimensions of the input and the kernel (or filter)? How many parameters are there in the kernel f?

Input: 6x6

Kernal: 3x3

of parameters are 9

Part3.2: What is the output activation map when you apply the convolutional operation using the filter f on the input X without padding?

The output will be:

$$C = ((n-f+2p)/s)+1$$

$$C = ((6-3+2 \times 0)/1)+1 = 4$$

Used online matrix calculator:

18	9	-4	-18
17	-3	-10	-12
11	-9	-17	2
9	-1	-15	16

Part3.3: What is the output when you apply a max-pooling operation on the output from the previous question?

Used online matrix calculator:

18	9
17	16